




Property Graph Databases




2018-10-01

Dr. Hamed Shariat Yazdi, Prof. Jens Lehmann

Motivation


Trend 1: Big Data



The amount of information that we generate and transfer has increased dramatically in the past few years.¹

¹Source: IDC's Digital Universe Study, sponsored by EMC, December 2012
<http://www.emc.com/collateral/analyst-reports/idc-the-digital-universe-in-2020.pdf>

Trend 2: Connectedness



Data is evolving to be more interlinked and connected. Hypertext has links, blogs have pingback, tagging groups all related data.

²Figure taken from [3]


Why Relational Databases are not Enough

(Provocative) claim: *Relational databases are not good at handling relationships!*

³Table taken from [4]

Why Relational Databases are not Enough

(Provocative) claim: *Relational databases are not good at handling relationships!*



In this social network database, it is easy to find people Bob considers his friends:


```
SELECT p1.Person
FROM Person p1 JOIN PersonFriend
  ON PersonFriend.FriendID = p1.ID
JOIN Person p2
  ON PersonFriend.PersonID = p2.ID
WHERE p2.Person = 'Bob'
```

Result:

³Table taken from [4]

Why Relational Databases are not Enough

(Provocative) claim: *Relational databases are not good at handling relationships!*



In this social network database, it is easy to find people Bob considers his friends:

```
SELECT p1.Person
FROM Person p1 JOIN PersonFriend
  ON PersonFriend.FriendID = p1.ID
JOIN Person p2
  ON PersonFriend.PersonID = p2.ID
WHERE p2.Person = 'Bob'
```

Result: Alice and Zach

³Table taken from [4]

Why Relational Databases are not Enough?

- ▷ Sad but true: friendship is not always symmetric!
- ▷ What if we want to know who considers Bob as friend?

```
SELECT p1.Person
FROM Person p1 JOIN PersonFriend
  ON PersonFriend.PersonID = p1.ID
JOIN Person p2
  ON PersonFriend.FriendID = p2.ID
WHERE p2.Person = 'Bob'
```

- ▷ User side: still easy to implement
- ▷ Database side: Database has to consider all the rows in the **PersonFriend** table \Rightarrow more expensive!

Why Relational Databases are not Enough?

What if we query deeper into the social network?

```
SELECT p1.Person AS PERSON, p2.Person AS FRIEND_OF_FRIEND
FROM PersonFriend pf1 JOIN Person p1
  ON pf1.PersonID = p1.ID
JOIN PersonFriend pf2
  ON pf2.PersonID = pf1.FriendID
JOIN Person p2
  ON pf2.FriendID = p2.ID
WHERE p1.Person = 'Alice' AND pf2.FriendID <> p1.ID
```

- ▷ Uses recursive joins \rightsquigarrow increase in syntactic, computational, and space complexity of the query

Why Relational Databases are not Enough?

Experimental results: Given any two persons chosen at random,
is there a path that connects them that is at most 5 relationships long?

Depth	RDBMS execution time (s)	Neo4j execution time (s)	Records returned
2	0.016	0.01	~ 2500
3	30.267	0.168	~ 110,000
4	1543.505	1.359	~ 600,000
5	Unfinished	2.132	~ 800,000

- ▷ Comparing relational store and a popular graph database Neo4j
- ▷ Social network consisting of 1,000,000 people
- ▷ Each person with approximately 50 friends

Conclusion: Graph databases outperform relational databases when we are dealing with connected data!

Relational Databases: What exactly is the Problem?

- ▷ Relationships exist only between tables → problematic for highly connected domains
- ▷ Relationship traversal can become very expensive
- ▷ Flat, disconnected data structures:
 - Data processing and relationship construction happens in the application layer
- ▷ Bound by a previously defined schema

Graph Databases using the Labeled Property Graph Model

- ▷ Explicit graph structure:
 - semantic dependencies between entities are made explicit
- ▷ New nodes and new relationships can be easily added into the database without having to migrate data or restructure the existing network
- ▷ Relationships correspond to paths:
 - querying the database = traversing the graph
→ this makes certain types of queries simpler
- ▷ Schema-free
- ▷ Index-free adjacency

- ▷ **Index-Free Adjacency:** each node “knows” (i.e. directly points to) its adjacent nodes
- ▷ No explicit index → each node acts as micro-index of nearby nodes → much cheaper than global indices
- ▷ Because of this, query times are less/not dependent of the size of the graph → they depend only on the part of the graph that has been searched
- ▷ Precomputes and stores bidirectional joins as relationships (i.e. no need to compute them)
- ▷ Enables bidirectional traversal of the database:
we can traverse both incoming and outgoing relationships

When to use:

- ▷ Complex, highly-connected data
- ▷ Dynamic systems with topology that is difficult to predict
- ▷ Dynamic requirements that change over time
- ▷ Cases where relationships in data are themselves meaningful

When not to use:

- ▷ Large offline batch analysis tasks of static data
- ▷ Simple, unconnected data structures
- ▷ For certain types of queries:
 - “Graph fishing” (no small set of start nodes)
 - “Bulk Scans” (generic queries not applying to any indexed subset)

Building blocks of a Labeled Property Graph:

▷ **Nodes**

- Can be tagged with one or more labels
- Can contain properties


▷ **Relationships:**

- Connect nodes and structure the graph
- Directed
- Always have a single name, a start node and an end node \Rightarrow no dangling relationships
- Can have properties

▷ **Properties** (= key-value pairs)

▷ **Labels:**


- Group nodes together




Properties are named values:

- ▷ Key is the name of the property
 - Always a string
- ▷ Values can be:
 - Numeric values
 - String values
 - Boolean values
 - Lists of any other type of value


- ▷ Used to assign types to nodes
- ▷ Groups nodes into sets:
all nodes with the same label belong to the same set
- ▷ Queries can be constrained to these sets instead of the whole graph \rightsquigarrow more efficient queries that are easier to write
- ▷ Each node has any number of labels (including none)






- ▷ Together with relationships, fundamental unit of property graph model
 \rightsquigarrow the simplest possible graph is a single node
- ▷ Are often used to represent entities
- ▷ Can have zero or more properties

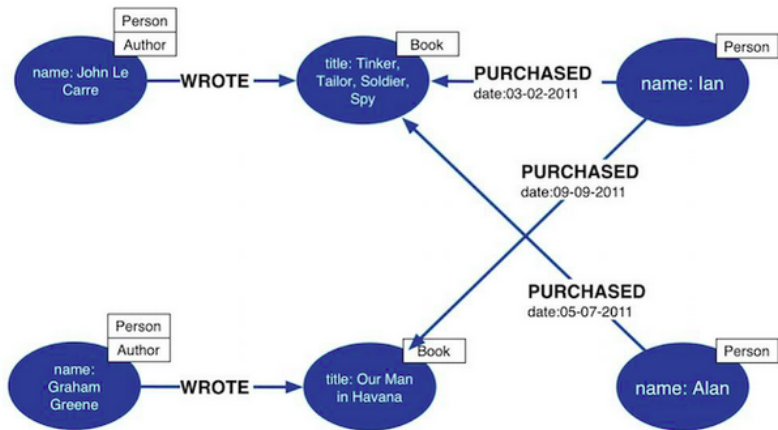
- ▷ Organize the nodes by connecting them
- ▷ Always connects a start node to the end node
- ▷ A key part of a graph database: allow us to find related data
- ▷ Always has a direction





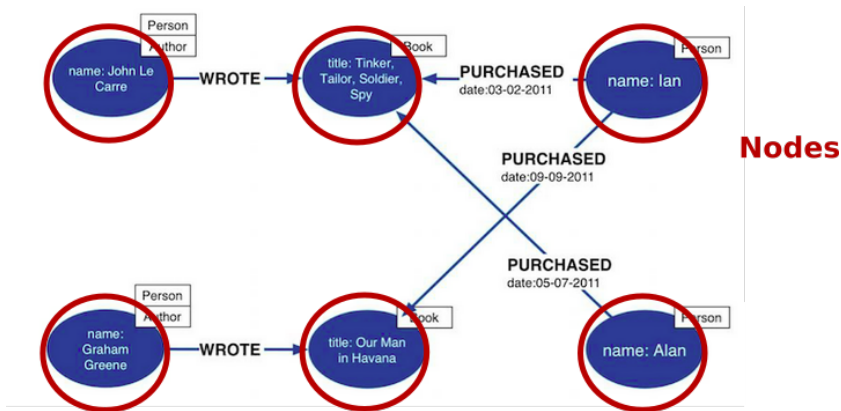
- ▷ One or more nodes with connecting relationships
- ▷ Typically is a result of a query or a traversal
- ▷ Length of a path = number of relationships on that path

Labeled Property Graph Model: Example

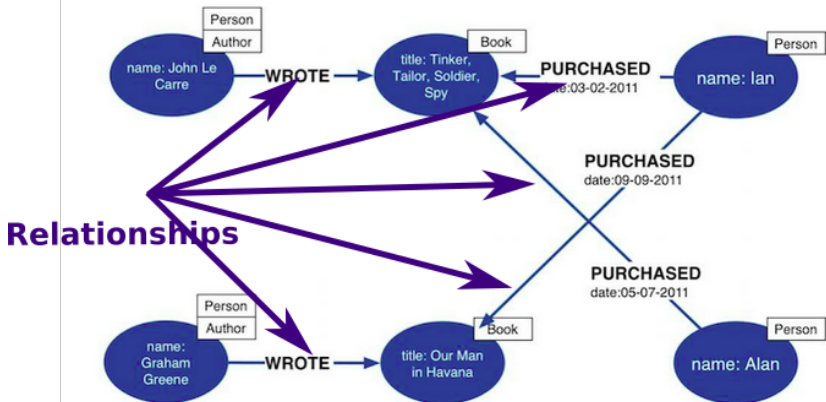


³Illustration taken from [1]

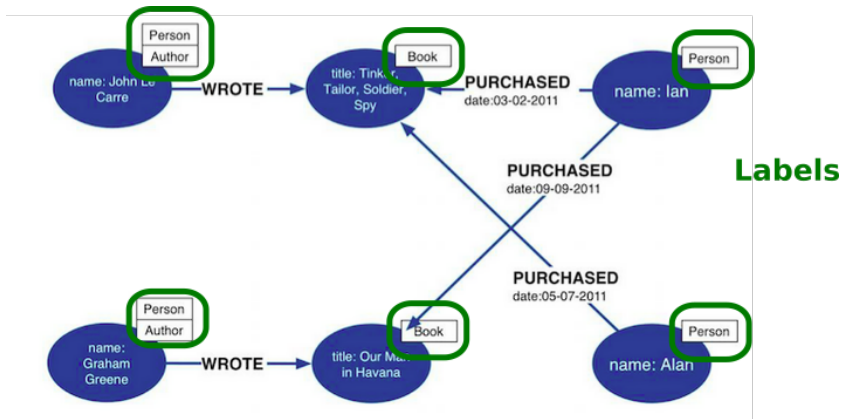
Labeled Property Graph Model: Example



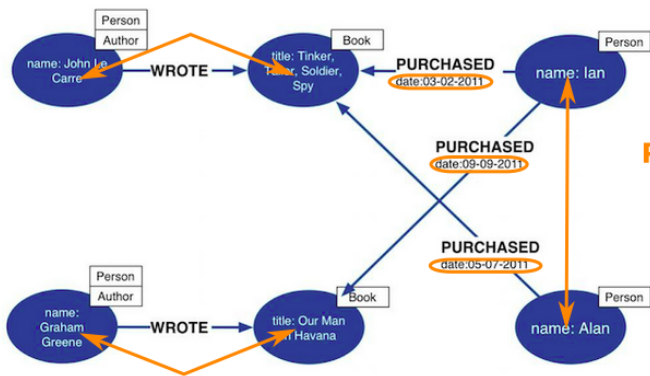
Labeled Property Graph Model: Example



Labeled Property Graph Model: Example



Labeled Property Graph Model: Example



Properties

Creating Graph Databases vs. Creating Relational Databases

Step 1: Acquire and improve understanding of data: a **whiteboard sketch** step

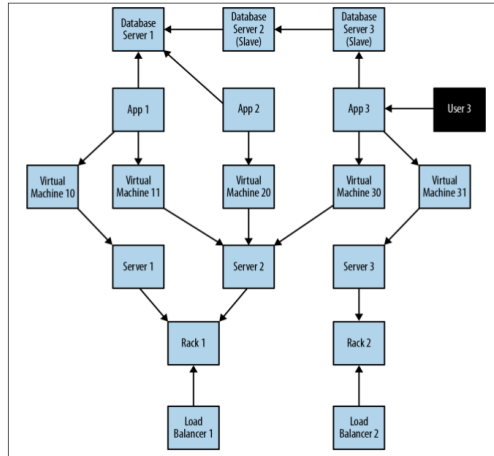


Figure 3-2. Simplified snapshot of application deployment within a data center

³Diagram taken from [4]

Step 2: Construct an entity-relationships (E-R) diagram




Figure 3-3. An entity-relationship diagram for the data center domain

Note the complexity of the model *before* any data has even been added

³E-R diagram taken from [4]

Step 3: Map the E-R diagram into tables and relations and normalize the data

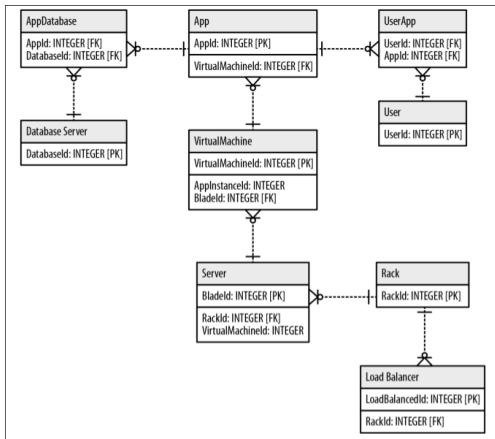


Figure 3-4. Tables and relationships for the data center domain

³Diagram taken from [4]

Creating a Graph Database

1. Create a whiteboard sketch
 2. Create an E-R diagram
 3. Map into tables and relations
- ← 2. Domain modeling

- ▷ What you sketch on the whiteboard = what you store in the database
- ▷ No normalization, denormalization or conversion to tables-relations structure necessary
- ▷ No conceptual-relational dissonance: the physical layout of the data is same as it is conceptualized
- ▷ Domain modeling = further graph modeling:
 - Makes sure that every node has the appropriate properties
 - Ensures that every node is in correct semantic context (i.e. add the relations you want to query)

Creating a Graph Database

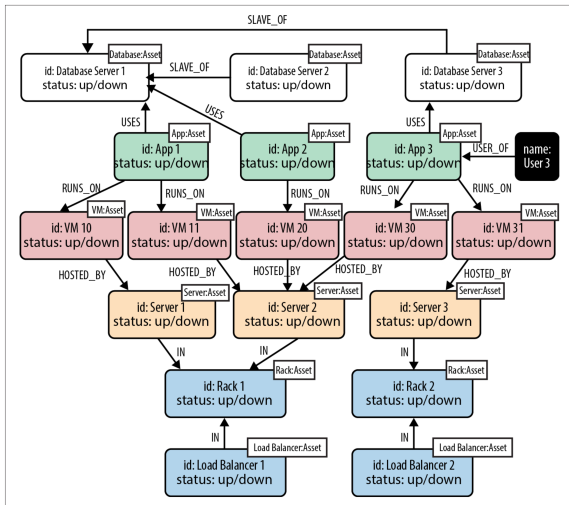


Figure 3-5. Example graph for the data center deployment scenario

³Image taken from [4]

- ▷ Ensure that later changes are driven by changes in application requirements and not by the need to mitigate bad design decisions
- ▷ There are two techniques to do this:
 - Check that the graph reads well
 - ▷ Pick a node
 - ▷ Follow relationships to other nodes, reading each node's role and each relationship's name as you go
 - ▷ *This should form sensible sentences*
 - *Design for queryability*
 - ▷ Understand end-users' goals: understand the use cases
 - ▷ Try to craft sample queries for the use cases

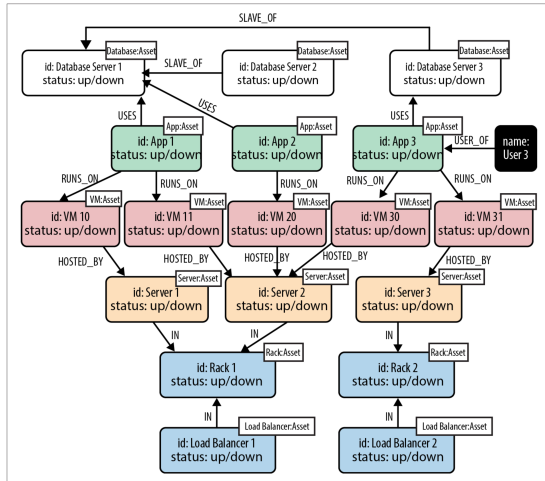


Figure 3-5. Example graph for the data center deployment scenario

“App 1 runs on VM 10, which is hosted by Server 1 in Rack 1.”

Neo4J - A Graph Database

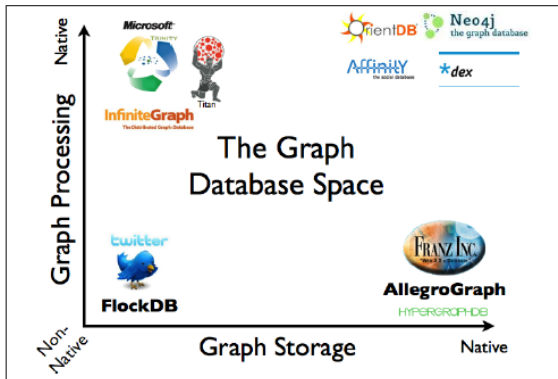


Figure 1-3. An overview of the graph database space

Source: O'Reilly Graph Databases [1]

- ▷ Native graph processing = index-free adjacency = traversal queries work well
- ▷ Native graph storage = storing data in graph shape (e.g. no RDBMS backend)



<https://neo4j.com/>

- ▷ Probably the most popular graph database today
- ▷ Based on a schema-free labeled property graph model
- ▷ Scales to billions of nodes, relationships and properties

Consists of:


- ▷ Nodes, Relationships, Properties, Labels, Paths, Traversal, and Schema (index and constraints)

- ▷ Open-source
- ▷ Extensive support and learning material
- ▷ Check out <https://neo4j.com/developer/> for further resources (including a sandbox for testing!)
- ▷ International events and meetup groups


(Neo4j)-[:LOVES]-(Developers)

Cypher


A pattern matching query language for graphs




A pattern matching query language for graphs



A pattern matching query language for graphs




Uses “ASCII art representation”




`() --> ()`

Directed relationship



(A) --> (B)

Undirected relationship



`(A) -- (B)`

Specific relationships




(A) -[:LOVES]-> (B)

Joined paths




(A) --> (B) --> (C)

Multiple paths



`(A) --> (B) --> (C), (A) --> (C)`
`(A) --> (B) --> (C) <-- (A)`



Cypher Patterns

```

(emil:Person {name:'Emil'})
<-[:KNOWS]- (jim:Person {name:'Jim'})
-[:KNOWS]-> (ian:Person {name:'Ian'})
-[:KNOWS]-> (emil)
  
```

- ▷ Specification by example: *draw* data we are looking for
- ▷ Used to define a pattern of nodes and relationships that we want to find

MATCH - example

```
MATCH (a:Person {name:'Jim'})-[:KNOWS]->(b)-[:KNOWS]->(c),  
      (a)-[:KNOWS]->(c)  
RETURN b, c
```

⇒ Reads: *“Find a node a with label person and name 'Jim'. Starting from a, find a neighbour node b via relation "KNOWS". Then find a neighbour node c of both a and b via relation "KNOWS".*

⇒ Short: Find mutual friends of Jim.

MATCH - example

```
MATCH (a:Person {name:'Jim'})-[:KNOWS]->()-[:KNOWS]->(c),  
      (a)-[:KNOWS]->(c)  
RETURN c
```

Same as before, but we are not interested in *b* this time.

- ▷ Cypher also provides various options to process the returned results
- ▷ They include options to aggregate, order, filter, and limit the returned data
- ▷ Example: `count(...)` allows us to return only the number of matched instances

RETURN Options - example

```
MATCH (theater:Venue {name:'Theatre Royal'},  
      (writer:Author {lastname:'Shakespeare'}),  
      (theater) <-[:VENUE]- (:Performance) -[:PLAY_OF]-> (writer)  
RETURN theater.city
```

Query: Cities with Shakespeare performances in theaters named “Theatre Royal”

RETURN Options - example

```
MATCH (theater:Venue {name:'Theatre Royal'}),  
      (writer:Author {lastname:'Shakespeare'}),  
      (theater) <-[:VENUE]- (:Performance) -[p:PLAY_OF]-> (writer)  
RETURN theater.city AS city, count(p) AS play_count
```

Query: Shakespeare performances in theaters named “Theatre Royal” counting the number of plays

Note: identifiers can also be attached to relations

RETURN Options - example

```
MATCH (theater:Venue {name:'Theatre Royal'}),  
      (writer:Author {lastname:'Shakespeare'}),  
      (theater) <-[:VENUE]- (:Performance) -[p:PLAY_OF]-> (writer)  
RETURN theater.city AS city, count(p) AS play_count  
ORDER BY play_count DESC
```

Query: Shakespeare performances in theaters named “Theatre Royal” ordered by the number of plays.

Note: assign/rename variable names with AS to use them in ORDER BY clause

RETURN Options - example

```
MATCH (theater:Venue {name:'Theatre Royal'}),  
      (writer:Author {lastname:'Shakespeare'}),  
      (theater) <-[:VENUE]- (:Performance) -[p:PLAY_OF]-> (writer)  
RETURN theater.city AS city, count(p) AS play_count  
ORDER BY play_count DESC  
LIMIT 1
```

Query: “Theatre Royal” with most Shakespeare plays

- ▷ WHERE constrains graph matches by one/more of the following constraints:
- presence/absence of certain paths in the matched subgraphs
 - certain labels for nodes
 - certain names for relationships
 - presence/absence of specific properties for matched nodes/relationships
 - specific values for properties of matched nodes/relationships
 - satisfaction of other constraints
e.g. those performances must have occurred before a certain date

For example, we can query specifically for Shakespeare plays that were written *after* 1608 (Shakespeare's final period):

WHERE - example

```
MATCH (bard:Author {lastname:'Shakespeare'}),  
      (play) <-[w:WROTE_PLAY]- (bard)  
WHERE w.year > 1608  
RETURN DISTINCT play.title AS play
```

Cypher supports a variety of clauses:

- ▷ **CREATE** and **CREATE UNIQUE**
- ▷ **DELETE**
- ▷ **SET**
- ▷ **FOREACH**
- ▷ **UNION**
- ▷ **WITH**

Remember our previous example!

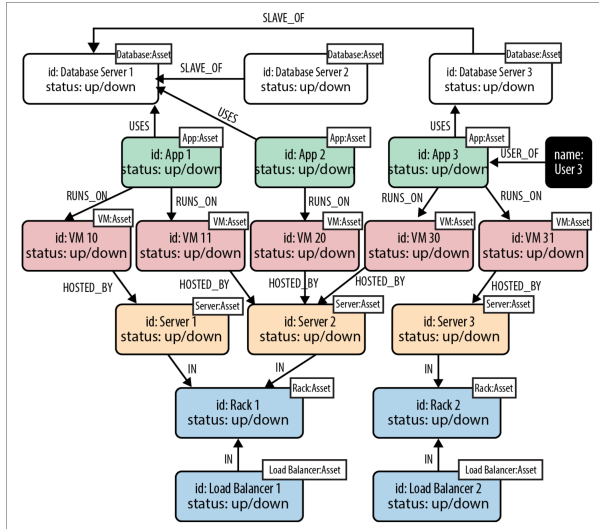


Figure 3-5. Example graph for the data center deployment scenario

Remember the design for queryability design goal!

Goal: Design a query to find the cause behind an unresponsive application or service in our example graph.

Example Query

```
MATCH (user:User)-[*1..5]-(asset:Asset)
WHERE user.name = 'User3' AND asset.status = 'down'
RETURN DISTINCT asset
```

The sample query we would need to define is:

Example Query

```
MATCH (user:User)-[*1..5]-(asset:Asset)
WHERE user.name = 'User3' AND asset.status = 'down'
RETURN DISTINCT asset
```

The sample query we would need to define is:

Example Query

```
MATCH (user:User)-[*1..5]-(asset:Asset)
WHERE user.name = 'User3' AND asset.status = 'down'
RETURN DISTINCT asset
```

- ▷ Describes a variable length path between one and five relationships long
- ▷ There is no colon or relationship name between the square brackets \rightsquigarrow the relationships are unnamed
- ▷ There are no arrow-tips \rightsquigarrow relationships are undirected

The sample query we would need to define is:

Example Query

```
MATCH (user:User)-[*1..5]-(asset:Asset)
WHERE user.name = 'User3' AND asset.status = 'down'
RETURN DISTINCT asset
```


The sample query we would need to define is:

Example Query

```
MATCH (user:User)-[*1..5]-(asset:Asset)
WHERE user.name = 'User3' AND asset.status = 'down'
RETURN DISTINCT asset
```

- ▷ We start with the user who reported a problem
- ▷ We add asset nodes that have a status property with a value of 'down'
- ▷ Nodes which do not have a status property will not be added to the results

The sample query we would need to define is:

Example Query

```
MATCH (user:User)-[*1..5]-(asset:Asset)
WHERE user.name = 'User3' AND asset.status = 'down'
RETURN DISTINCT asset
```

The sample query we would need to define is:

Example Query

```
MATCH (user:User)-[*1..5]-(asset:Asset)
WHERE user.name = 'User3' AND asset.status = 'down'
RETURN DISTINCT asset
```

- ▷ Ensures that unique assets are returned in the results, no matter how many times they are matched

Example 1

```
MATCH (p:Product {productName: "Chocolate" })  
RETURN p.productName, p.unitPrice
```

Example 1

```
MATCH (p:Product {productName: "Chocolate" })  
RETURN p.productName, p.unitPrice
```

- ▷ Gets chocolates and their price
- ▷ Alternative:

Example 1 Alternative

```
MATCH (p:Product)  
WHERE p.productName = "Chocolate"  
RETURN p.productName, p.unitPrice
```

Example 2

```
MATCH (p:Product)
RETURN p.productName, p.unitPrice
ORDER BY p.unitPrice DESC
LIMIT 10
```

Example 2

```
MATCH (p:Product)
RETURN p.productName, p.unitPrice
ORDER BY p.unitPrice DESC
LIMIT 10
```

- ▷ Returns only a subset of attributes, in this case: ProductName and UnitPrice
- ▷ Orders by price
- ▷ Returns 10 most expensive items


Example 3

```
MATCH  (p:Product {productName:'Chocolade'})  
       <-[:PRODUCT]- (:Order)  
       <-[:PURCHASED]- (c:Customer)  
RETURN DISTINCT c.name
```


Example 3

```
MATCH  (p:Product {productName:'Chocolade'})  
       <-[:PRODUCT]- (:Order)  
       <-[:PURCHASED]- (c:Customer)  
RETURN DISTINCT c.name
```

▷ Names of everyone who bought *Chocolade*




Example 4

MATCH (bob:User {username:'Bob'})-[:SENT]->(email)-[:CC]->(alias),
(alias)-[:ALIAS_OF]->(bob)

RETURN email.id

³Example and figure taken from [4]




Example 4

MATCH (bob:User {username:'Bob'})-[:SENT]->(email)-[:CC]->(alias),
 (alias)-[:ALIAS_OF]->(bob)
RETURN email.id

- ▷ Returns all emails that Bob has sent where he's CC'd one of his own aliases
- ▷ Returns 1 result: id: "1", content: "email content"


³Example and figure taken from [4]



Example 5

MATCH (email:Email {id:'11'}) < -[f:FORWARD_OF*]-(:Forward)
RETURN count(f)

³Example and figure taken from [4]




Example 5

MATCH (email:Email {id:'11'}) < -[f:FORWARD_OF*]-(:Forward)
RETURN count(f)

- ▷ This returns the number of times a particular email is forwarded
- ▷ The answer is 2 (count the number of FORWARD_OF relationships bound to f)

³Example and figure taken from [4]

Serialisation



```

CREATE (:Person {name:'Ian'})-[:EMPLOYMENT]->
  (employment:Job {start_date:'2011-01-05'})
-[:EMPLOYER]-> (:Company {name:'Neo'}),
  (employment)-[:ROLE]-> (:Role {name:'engineer'})
  
```

³Image taken from [4]

- ▷ GEOFF = **Graph Export Object File Format**
- ▷ Is a text representation of a graph
- ▷ Based on Cypher
- ▷ A GEOFF document consists of a one or more subgraphs, each of which contains one or more paths
- ▷ Properties are in JSON syntax

```
(alice {"name":"Alice"})  
(bob {"name":"Bob"})  
(carol {"name":"Carol"})  
(alice)<-[:KNOWS]->(bob)<-[:KNOWS]->(carol)<-[:KNOWS]->(alice)
```


- ▷ XML-based file format for graphs
- ▷ Common format for exchanging graph structure data
- ▷ Generic (not limited to graph databases)

```
<?xml version="1.0" encoding="UTF-8"?>
<graphml xmlns="http://graphml.graphdrawing.org/xmlns"
[...]
```

```
  <graph id="G" edgedefault="undirected">
    <node id="n0">
      <data key="d0">green</data>
    </node>
    <node id="n1"/>
    <edge id="e0" source="n0" target="n1">
      <data key="d1">1.0</data>
    </edge>
  </graph>
</graphml>
```

- ▷ Essentially no standardisation in property graph community
- ▷ Serialisation is less important than in semantic technologies:
 - Not much emphasis on data exchange
 - Data publishing not commonly considered (in contrast to Linked Data)
 - Embedding in other formats usually not considered (in contrast to RDFa)

Converting and Comparing the Graph Models

Distinguish between two types of RDF triples:

- ▷ *attribute triples*
= triples whose object is a literal
- ▷ *relationship triples*
= triples whose object is an IRI or a blank node

The transformation:

- ▷ Every relationship triple \Rightarrow an edge
- ▷ Every attribute triple \Rightarrow a property of the vertex for the subject of that triple
- ▷ IRI is preserved via its own property

Optional additional conversion:

Triples with predicate `rdf:type` can be used to assign labels to nodes (as the meaning of a type assignment resembles that of a label).


Simple transformation example:



```

ex:alice foaf:knows ex:bob .
ex:alice foaf:name "Alice" .
ex:bob foaf:name "Bob" .
  
```

Reification transformation example:



```
ex:alice foaf:name "Alice" .
ex:bob foaf:name "Bob" .
<<ex:alice foaf:knows ex:bob>> ex:certainty 0.5 .
<<ex:bob foaf:age 23>> ex:certainty 0.9 .
```

- ▷ We can not have labels for predicates in the RDF format (e.g. `ex:certainty=0.5` as above)
- ▷ A solution is to use Turtle* syntax (an extension of the RDF Turtle format)
- ▷ Turtle* embeds RDF triples into other RDF triples by enclosing the embedded triple in `<<` and `>>` and use it as subject/object
- ▷ However the transformation of the last sentence is not possible
See: <https://arxiv.org/pdf/1409.3288v2.pdf> for more details

- ▷ Relationships \Rightarrow a (relationship) RDF triple
- ▷ Node properties (incl. their labels) \Rightarrow an (attribute) RDF triple
- ▷ Relationship properties \Rightarrow metadata triple whose subject is the triple for the corresponding edge
- ▷ Patterns for generating IRIs that denote edge labels and properties can be chosen freely

Characteristics	Relational	RDF	PGM
Standardised	yes	yes	no
Traversal performance	—	~	+
Large analytical queries	+	~	—
Query language	SQL	SPARQL	Cypher & more
Data Publication & Dereferencing	—	+	—
Global Identifiers & Cross dataset fusion	—	+	—

Legend:

- + : good performance
- ~ : medium performance
- : low performance

Covered in the lecture:

- ▷ Motivation for Model Graph Databases (Big Data, Connectivity)
- ▷ Comparing Relational and Property Graph Databases
- ▷ Labeled Property Graph Model
- ▷ Cypher Query Language
- ▷ “Unifying” RDF and Property Graphs
- ▷ Not covered: Hypergraph model (edges can have more than two vertices)



Neo4j.

<https://neo4j.com/developer/graph-database/#property-graph>.



O. Hartig.

Reconciliation of rdf* and property graphs.

arXiv preprint arXiv:1409.3288, 2014.



T. Ivarsson.

Graph database and neo4j.

[http:](http://www.slideshare.net/thobe/nosqleu-graph-databases-and-neo4j)

[//www.slideshare.net/thobe/nosqleu-graph-databases-and-neo4j](http://www.slideshare.net/thobe/nosqleu-graph-databases-and-neo4j).



I. Robinson, J. Webber, and E. Eifrem.

Graph Databases: New Opportunities for Connected Data.

" O'Reilly Media, Inc.", 2nd edition, 2015.